

Neural Lagrangian Mesh: A Generative Model for Manifold Meshes of Arbitrary Genus

Adrish Dey¹ Edward Chien¹

¹Boston University

BOSTON
UNIVERSITY

Introduction

We explore the use of Lagrangian mesh representation to build a conditional generative model for manifold meshes of arbitrary genus, conditioned on shape descriptors (currently images).

We introduce a conditional flow-matching framework to generate the particles and a neural convex solver based on input features. Next, we treat the object as a collection of “volumetric parcels,” defined as convex polytopes within a bounded continuous domain Ω , allowing us to leverage OT machinery. We also explore a Neural Convex solver for faster solutions. Finally, we employ skin-surface meshing (Edelsbrunner ’99) to obtain the manifold mesh.

Flow Matching for Particle Generation

Given a conditional feature vector z_j , we aim to generate a set of samples $\{y_i\}_{[N]}$ of the target shape \mathcal{O}_j . To achieve this, we leverage Flow Matching to parameterize a conditional integration map $\phi(x|z_j) \stackrel{\text{def}}{=} \int_0^1 v_\theta(x, z_j, t) dt$ acting as a pushforward between $\mathcal{N}(0, 1) \rightarrow \nu$. Here, $v_\theta(x, z_j, t)$ is the neural network parameterizing the vector field, and $\mathcal{N}(0, 1)$ denotes the isotropic Gaussian. Next, we design the conditional vector field $u(x|z_j)$. We follow OT-CFM by (Tong et al., 2023) to build an Optimal Transport gradient field as follows:

$$u(x|z_j) = x_1^* - x_0^*$$

where $(x_0^*, x_1^*) \sim \Pi(\mathbf{x}_0, \mathbf{x}_1)$ and $\mathbf{x}_0 \sim \mathcal{N}(0, 1)$, $\mathbf{x}_1 \sim P(\mathcal{O}_j)$. $\Pi(\mathbf{x}_0, \mathbf{x}_1)$ denotes the Optimal Transport coupling between the sample minibatches \mathbf{x}_0 and \mathbf{x}_1 .

The final objective function is:

$$\mathcal{L}_\theta^{\text{CFM}} = \|v_\theta(x, z_j, t) - u(x|z_j)\|^2$$

Semi-Discrete Optimal Transport

Consider the continuous background $\Omega := [0, 1]^3$. Let $\{y_i\}_{[N]}$ be the set of N generated particles. We equip the set of points $\{y_i\}_{[N]}$ with Dirac masses $\{\nu_i\}_{[N]}$ such that $\sum_i \nu_i < \|\Omega\|$. From the Lagrangian perspective, each point denotes the location of the volumetric parcels $\{V_i\}_{[N]}$ that constitute the mesh. Our goal is to find the volume parameter $\{\psi_i\}_{[N]}$ of the particles.

Let $\psi \stackrel{\text{def}}{=} \{\psi_i\}_{[N]}$. Following (Levy, 2021), let us define V_i^ψ as the “volume parcel” with volume parameter ψ_i and the background “air” as a continuum volume of remaining mass $\|\Omega\| - \sum_i \nu_i$.

Note: Lag_i^ψ denotes the i^{th} cell of the induced Laguerre tessellation parameterized by ψ . This gives rise to a **convex** optimization objective as follows:

$$\arg \max_{\psi \in \mathbb{R}_+^N} K(\psi) = \sum_{i \in [N]} \int_{V_i^\psi} \|x - y_i\| d\mu + \sum_{i \in [N]} y_i \nu_i$$

where,

$$V_i^\psi = Lag_i^\psi \cap B(y_i, \sqrt{\psi_i})$$

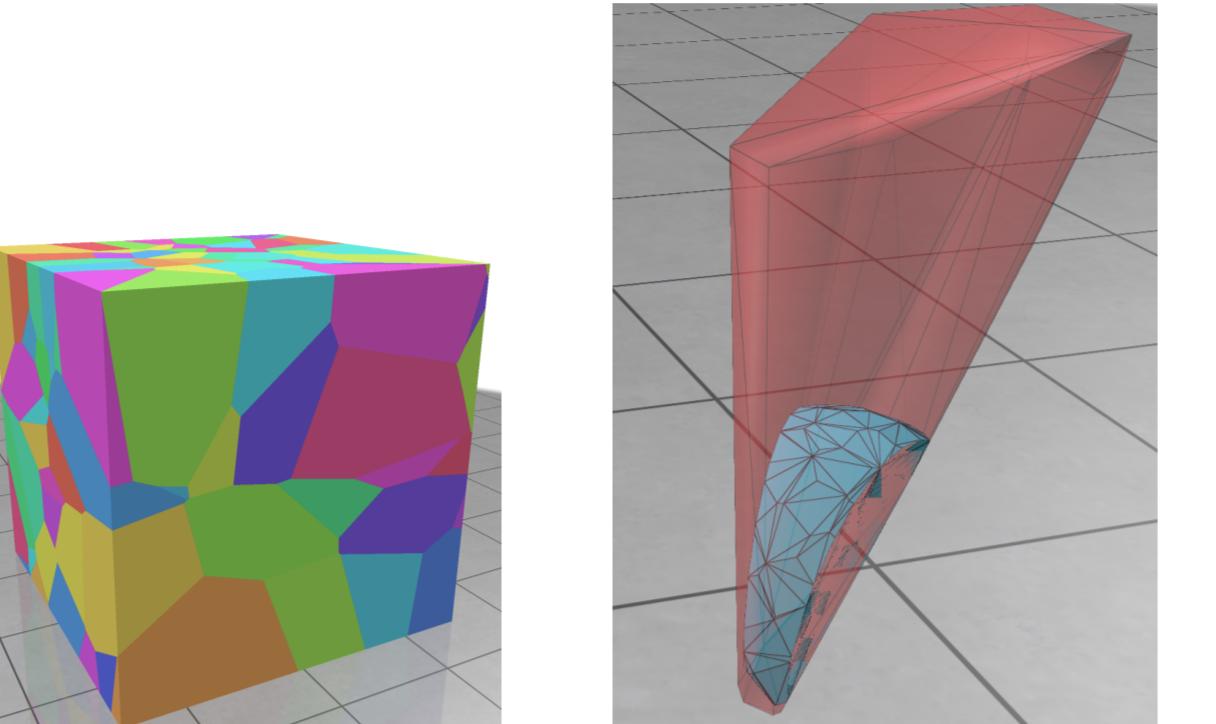


Figure 2. Left: Laguerre Tessellation of Domain. Right: (Blue) V_i^ψ cell. (Red) Laguerre Cell Lag_i^ψ

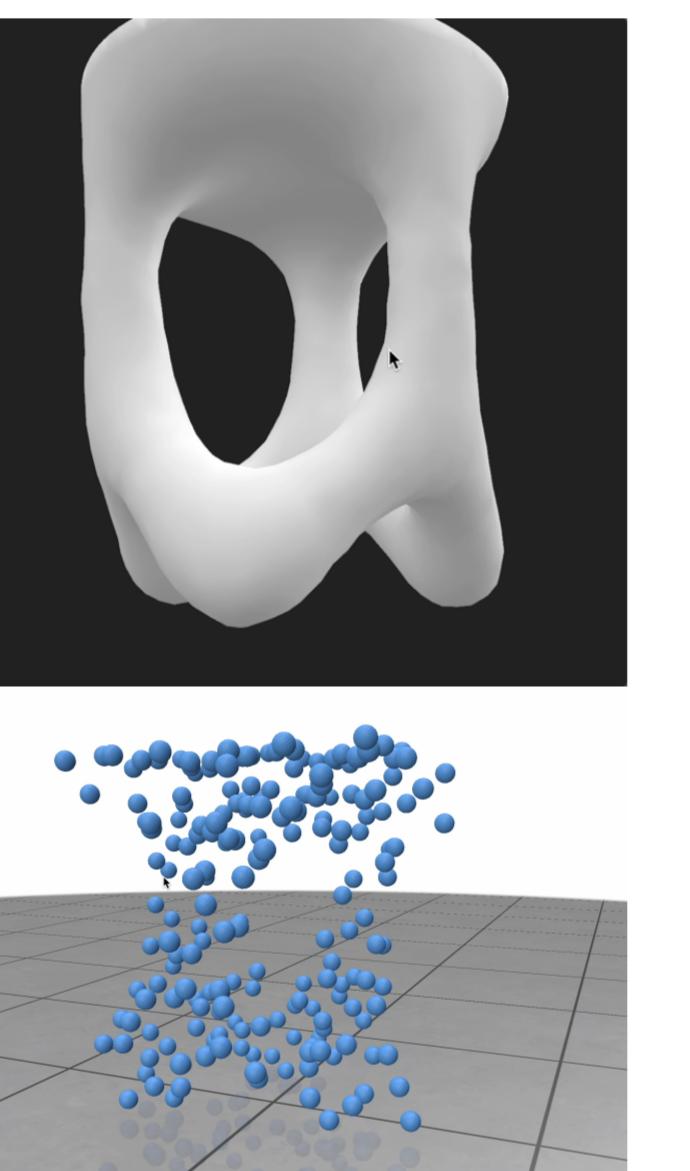


Figure 1. Top: Conditioning Image. Bottom: Generated Point Cloud

```

1:  $\psi_i \leftarrow 0$ 
2: while  $\left\| \frac{\partial K}{\partial \psi_i} \right\|^2 \geq \epsilon$  do
3:   Compute  $V_i := Lag_i^\psi \cap B(y_i, \sqrt{\psi_i})$ 
4:    $\frac{\partial K}{\partial \psi_i} = \nu_i - |V_i^\psi|$ 
5:    $\psi_i \leftarrow \psi_i + \lambda \cdot \frac{\partial K}{\partial \psi_i}$ 
6: end while

```

Algorithm 1. Gradient Ascent to find ψ_i

Computing V_i and $|V_i|$

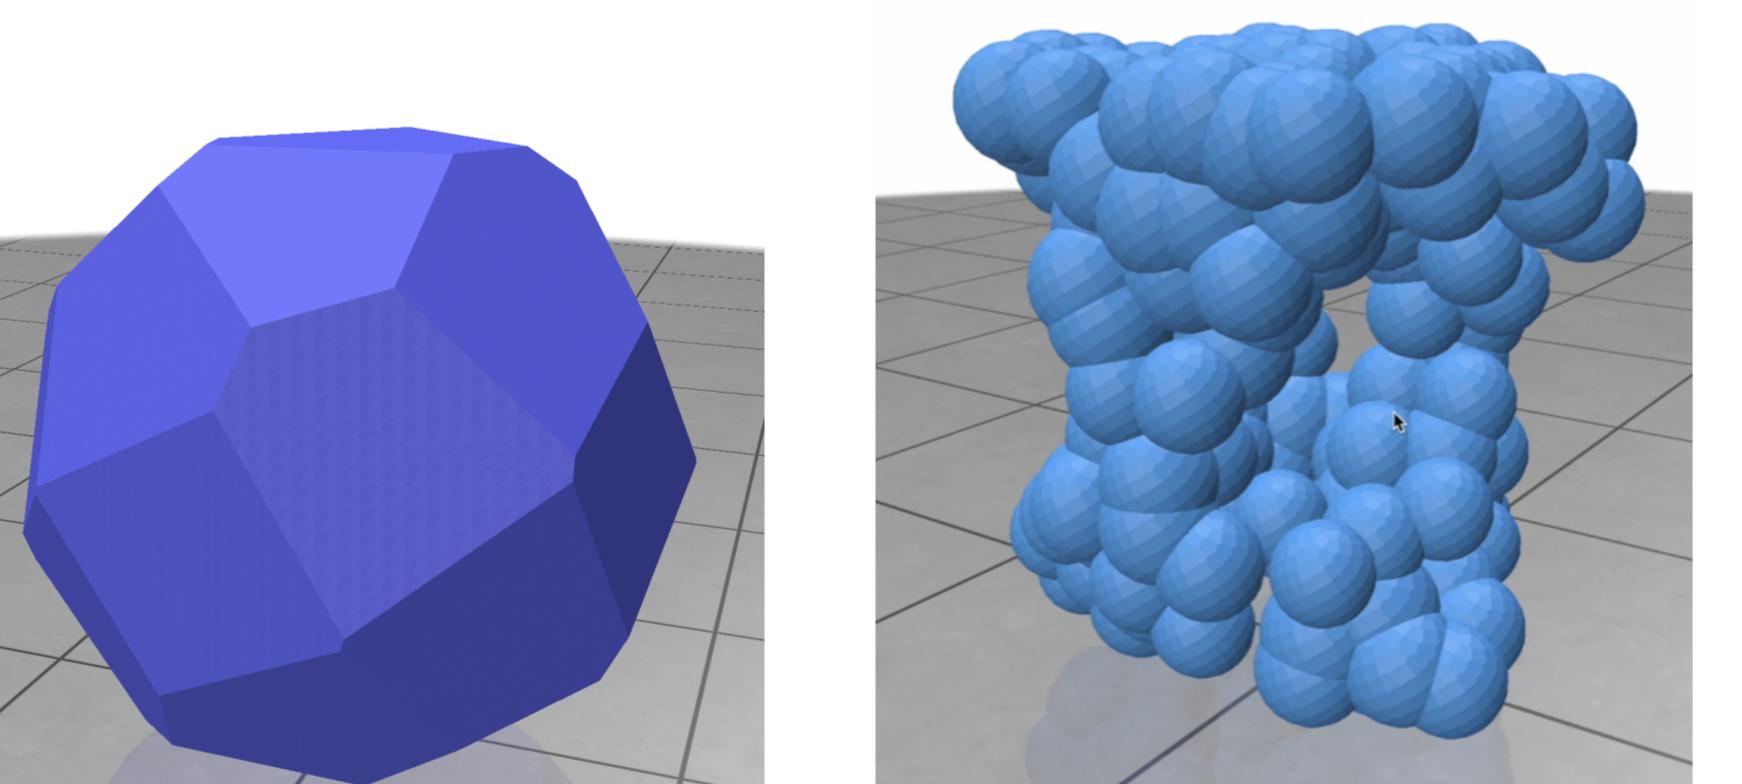


Figure 3. Left: Approximating $B(y_i, \sqrt{\psi_i})$ with a convex polytope of K -walls. Right: union of polytopes after optimization

For fast intersection of Lag_i^ψ with $B(y_i, \sqrt{\psi_i})$ we approximate each ball with a convex polytope $\hat{B}(y_i, \sqrt{\psi_i})$ with K walls. (NOTE: A large value of K is preferred for better approximation. We choose $K = 200$)

Polytope Intersection: To perform polytope intersection $V_i^\psi = Lag_i^\psi \cap \hat{B}(y_i, \sqrt{\psi_i})$, we first represent both Lag_i^ψ and $\hat{B}(y_i, \sqrt{\psi_i})$ using half spaces (i.e., H-Representation). Now, the desired polytope V_i^ψ is essentially the one which satisfies the conditions $\text{HalfSpace}(Lag_i^\psi) \cup \text{HalfSpace}(\hat{B})$. (Redundant walls are removed by using Chebychev’s radius).

Polytope Volume Computation: We first convert the H-representation of V_i^ψ to V-representation. The volume of the polytope is then calculated by tetrahedralizing the convex hull of V_i^ψ .

Gradient Flows

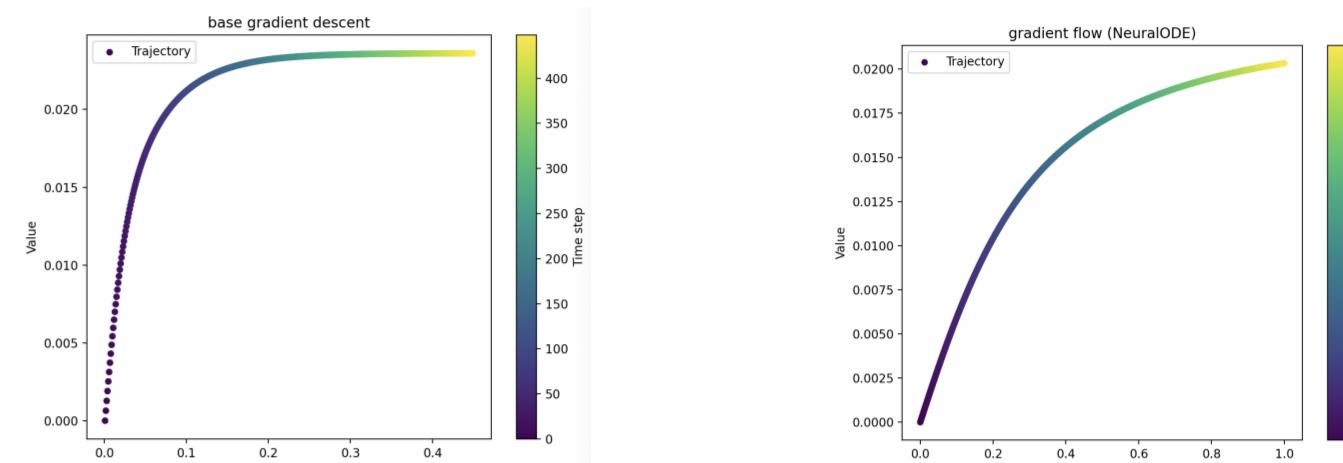


Figure 4. Gradient Flow Matching. Left: Value of potentials ψ_i at each discrete timestep t of gradient ascent. (Time taken: ~20min for 1000 steps) Right: Value of potentials obtained from learned continuous Gradient Flow. (Time taken: ~2 sec for the entire flow)

We are also interested in building a Neural Solver for the convex optimization in Algorithm 1, since solving this convex program (specifically computing the polytope V_i^ψ) is computationally expensive. To that end, we consider the gradient ascent algorithm as an Euler discretization of the gradient flow from $\psi_i = 0$ to ψ_i^{OPT} . Let T be the number of steps taken in the gradient ascent, and z_j the conditioning variable. We define a conditional integration map $\varphi_i(0|z_j) \stackrel{\text{def}}{=} \int_0^T \frac{\partial K}{\partial \psi_i} [\psi_i(t)] dt$ that solves $\psi_i^{OPT} = \varphi_i(0|z_j)$ for a given shape with descriptors z_j . Now, given a parametric function $\varrho_\theta(\psi_i(t), z_j, t)$, we formulate a flow matching algorithm that approximates the solution $\{\psi_i^{OPT}\}_{[N]}$. The conditional vector field is given as follows:

$$u(\psi_i(t)|z_j) = \frac{\partial K}{\partial \psi_i} [\psi_i(t)]$$

And the flow matching loss is given as follows:

$$\mathcal{L}_\theta^{\text{CFM}} = \|\varrho_\theta(\psi_i(t), z_j, t) - u(\psi_i(t)|z_j)\|^2$$

Skinsurface Meshing (Edelsbrunner'99, Kruithof et.al. '07)

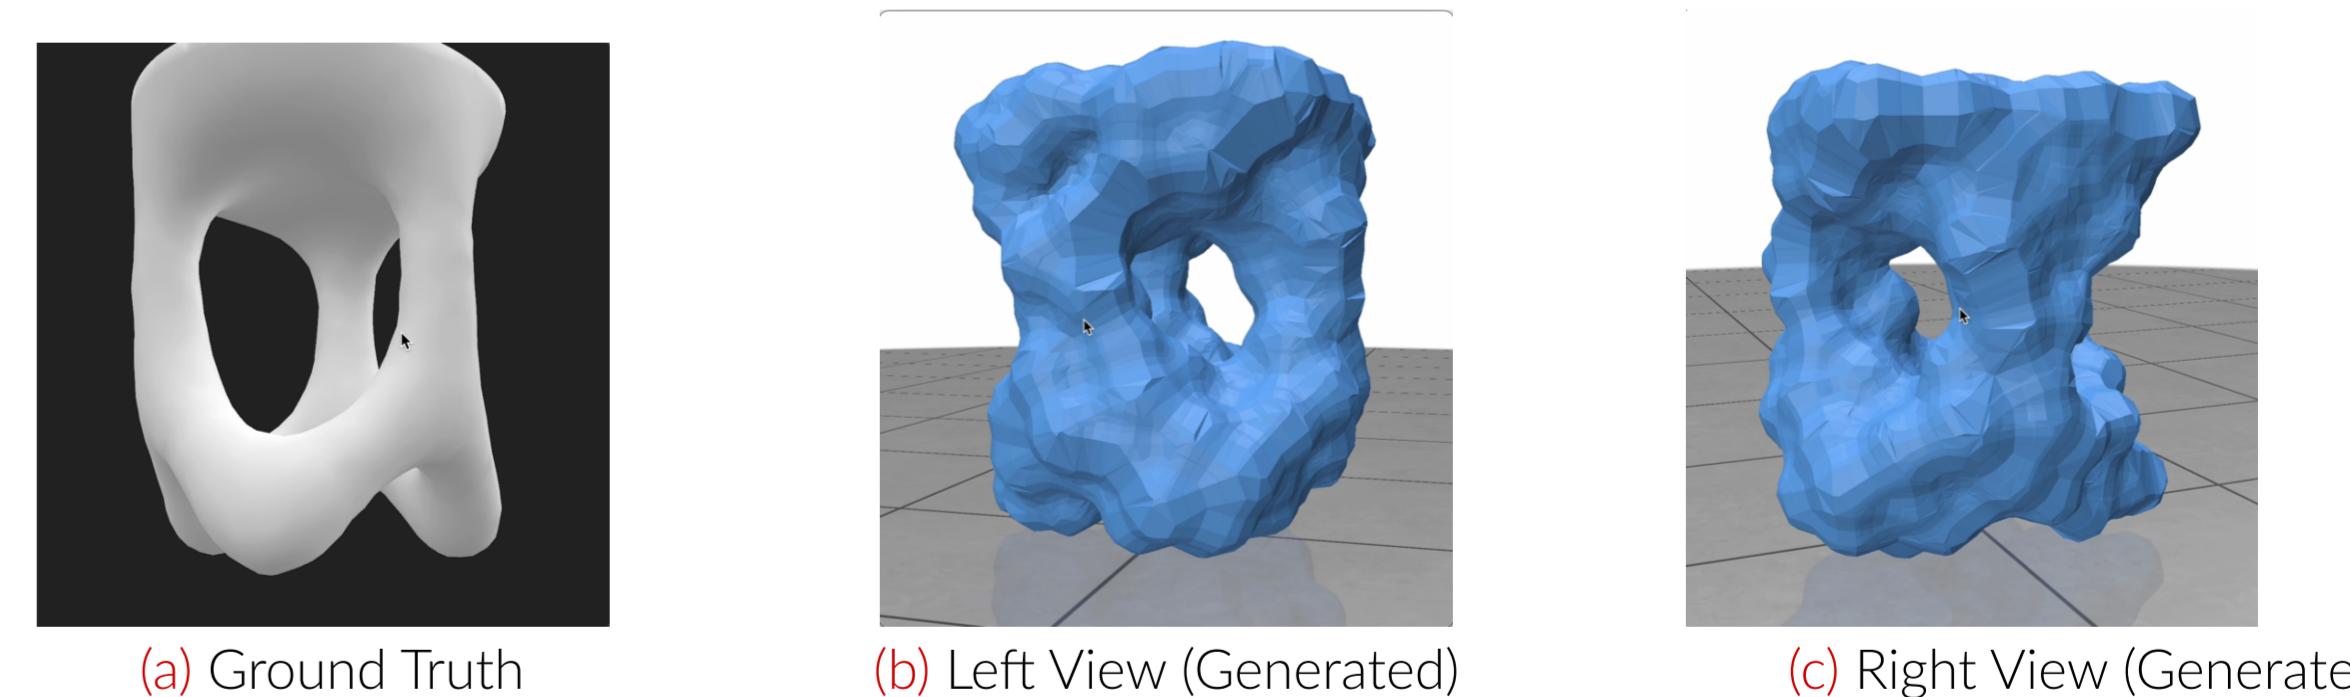


Figure 5. Skin surface meshing provides an efficient way to extract manifold surface mesh of arbitrary topology from a union of balls system.

Looking for Input

- What are the advantages of a Lagrangian representation? Does the flow-based formulation allow for any correspondence between various generated geometries?
- Can we combine these ideas with Gaussian splat representations? They are surface representations, while ours are volumetric.
- Can we combine these ideas with diffusion models for other generative purposes? Perhaps with branching to allow for adaptive density for more/less geometric detail.